

# 10 Efficient mesh models

Follows chapter 16 in RTR 4e

Steve Marschner

**CS5625** Spring 2019

# Basics of efficiency for meshes

## Use triangle or quad meshes

- general polygon meshes lead to too much complexity
- quad meshes are great for some applications but more constrained

## Use shared-vertex triangle meshes for GPU applications

- major memory/bandwidth savings over separate triangles
- if you get separate triangles, merge them in a pre-process

## Store most data at vertices

- there are ~half as many vertices as faces
- vertex data may be interpolated across faces
- in typical GPU mesh representation, vertices must be duplicated to create discontinuities

# More sophistication in mesh storage

## Optimizing vertex order

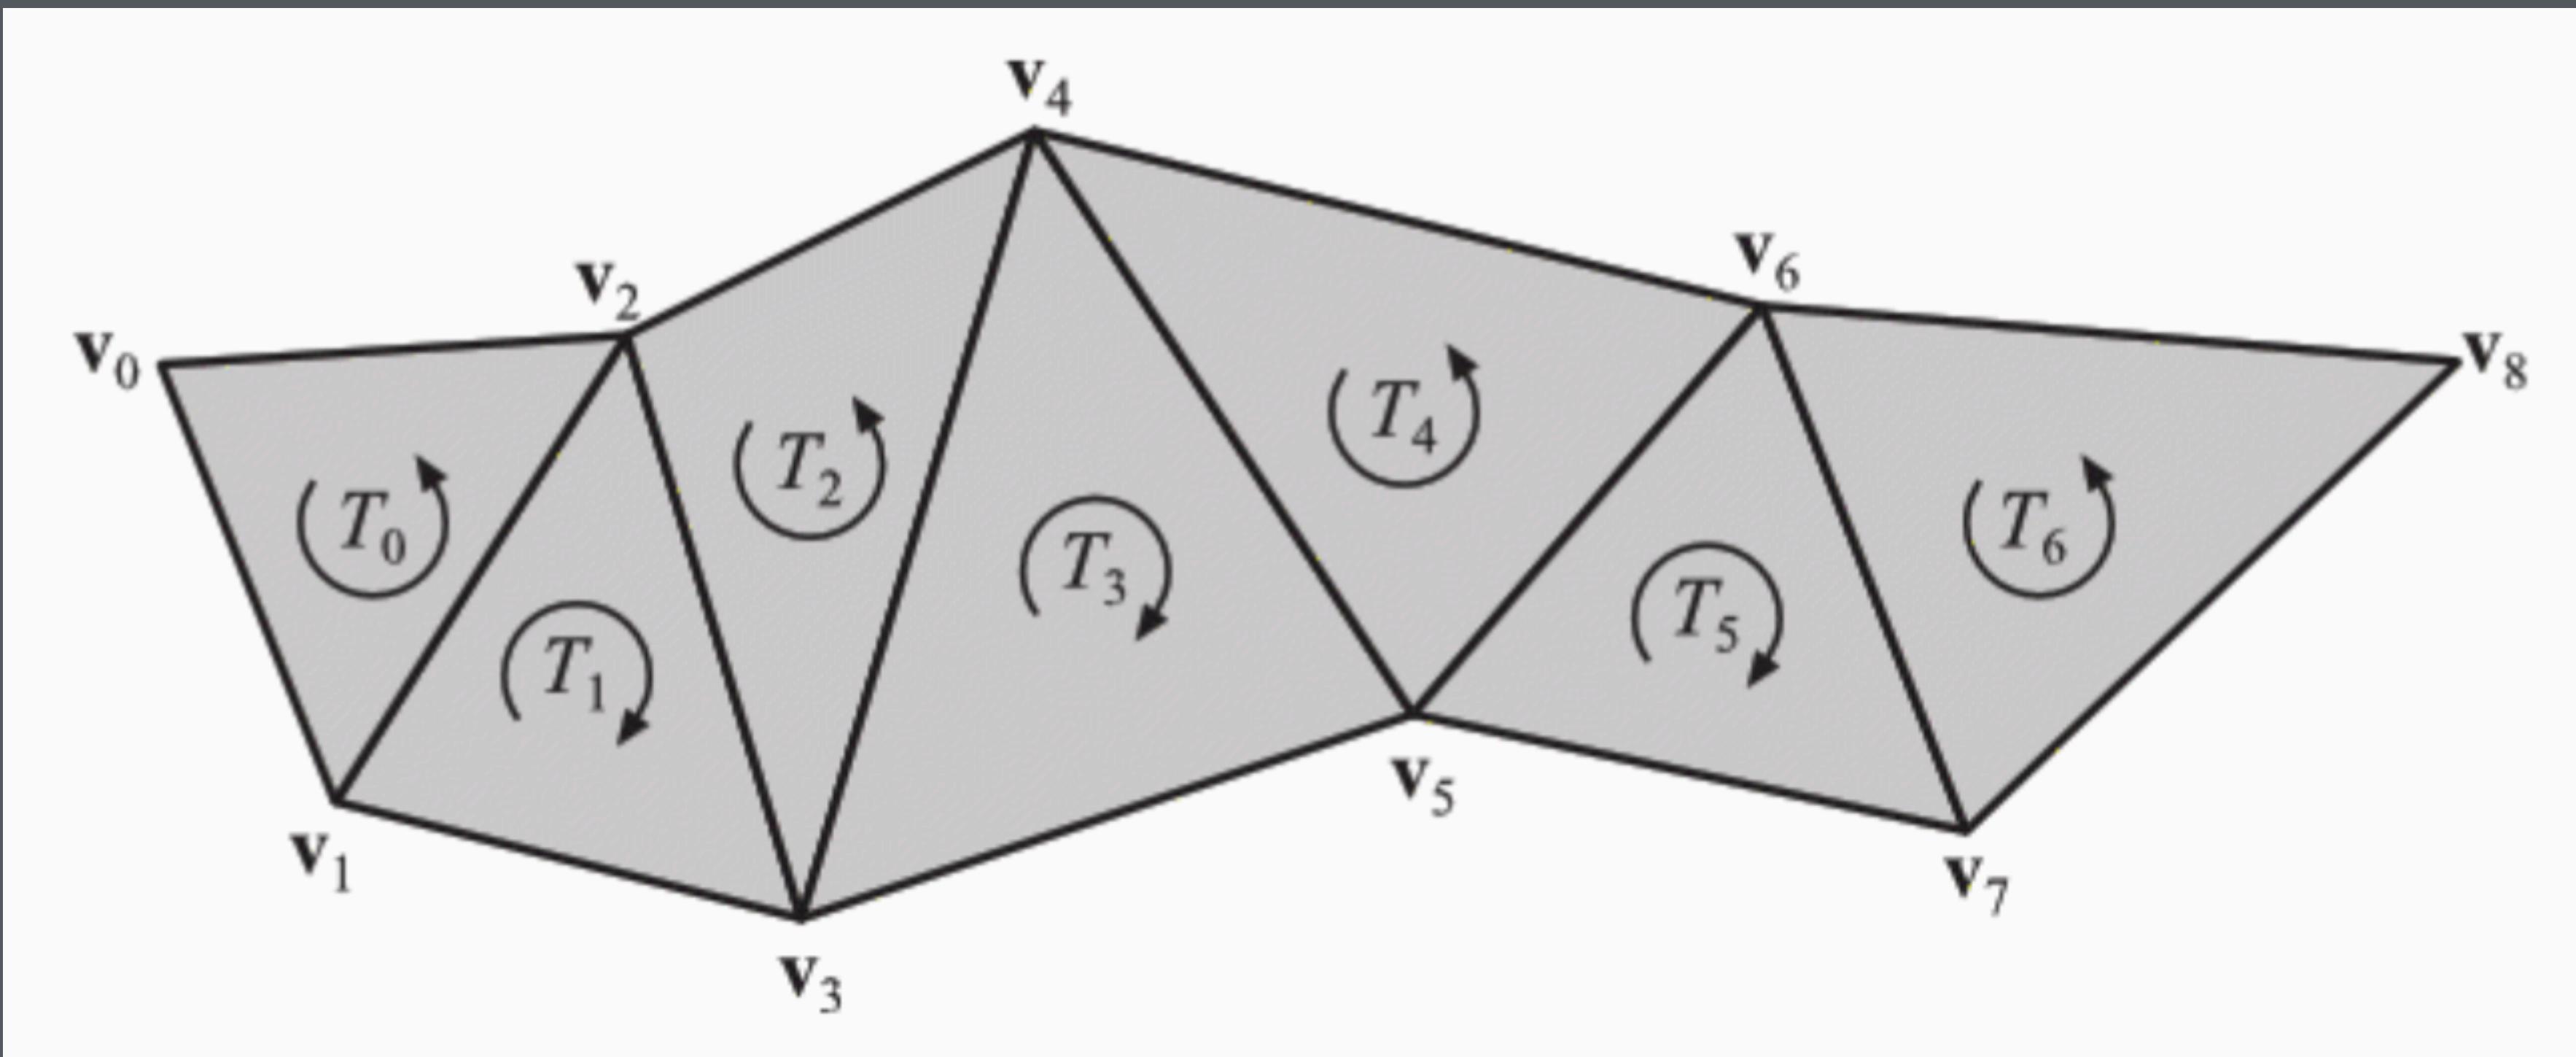
- strips and fans as examples (when per-frame bandwidth was the concern)
- modern systems don't use these but optimize for hit rate in vertex cache

## Reducing the number of triangles

- ultimately this is needed to save more time and space
- many *levels of detail* are useful
  - simpler meshes for faraway objects
  - simpler meshes for lower-resolution screens
  - simpler meshes for lower-performance hardware or networks

# Vertex cache and mesh ordering

**Triangle strips gain efficiency by caching the most recent two vertices**



- we are essentially using a FIFO cache policy with a size of 2
- cache miss rate approaches 1 miss / triangle

# Optimizing for larger caches

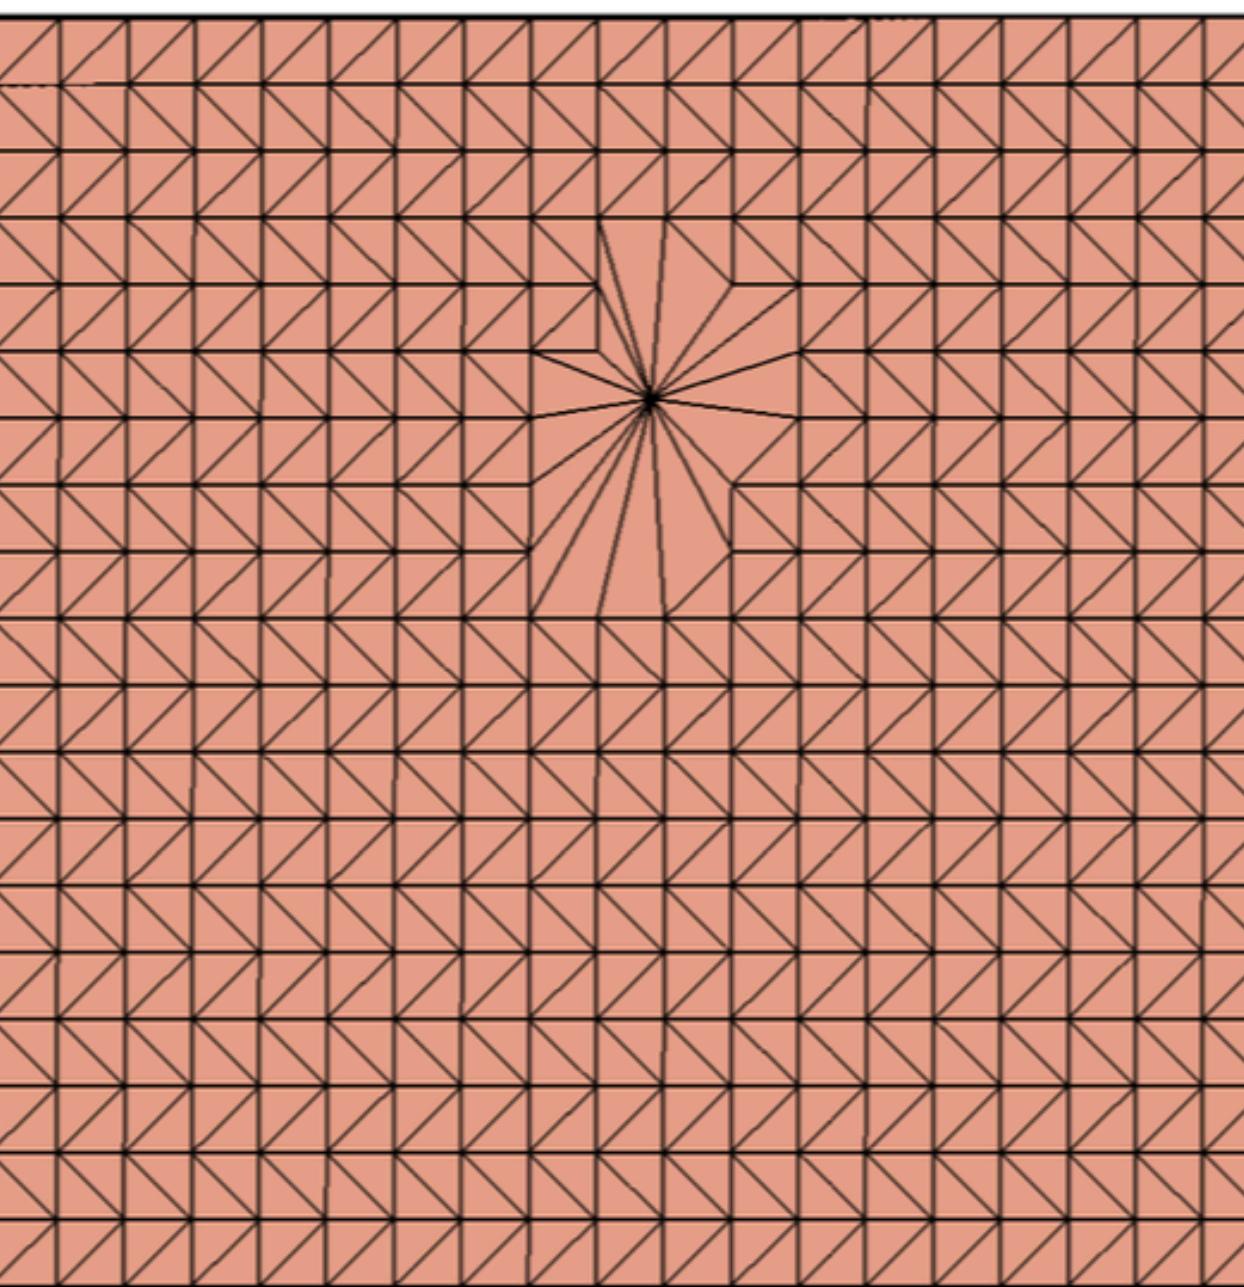
## **With indexed meshes, saving indices is less important**

- we store lots of data at vertices; ~6 indices is the least of our worries
- just putting meshes in triangle-strip order gives you the same vertex caching behavior (“transparent” vertex caching)

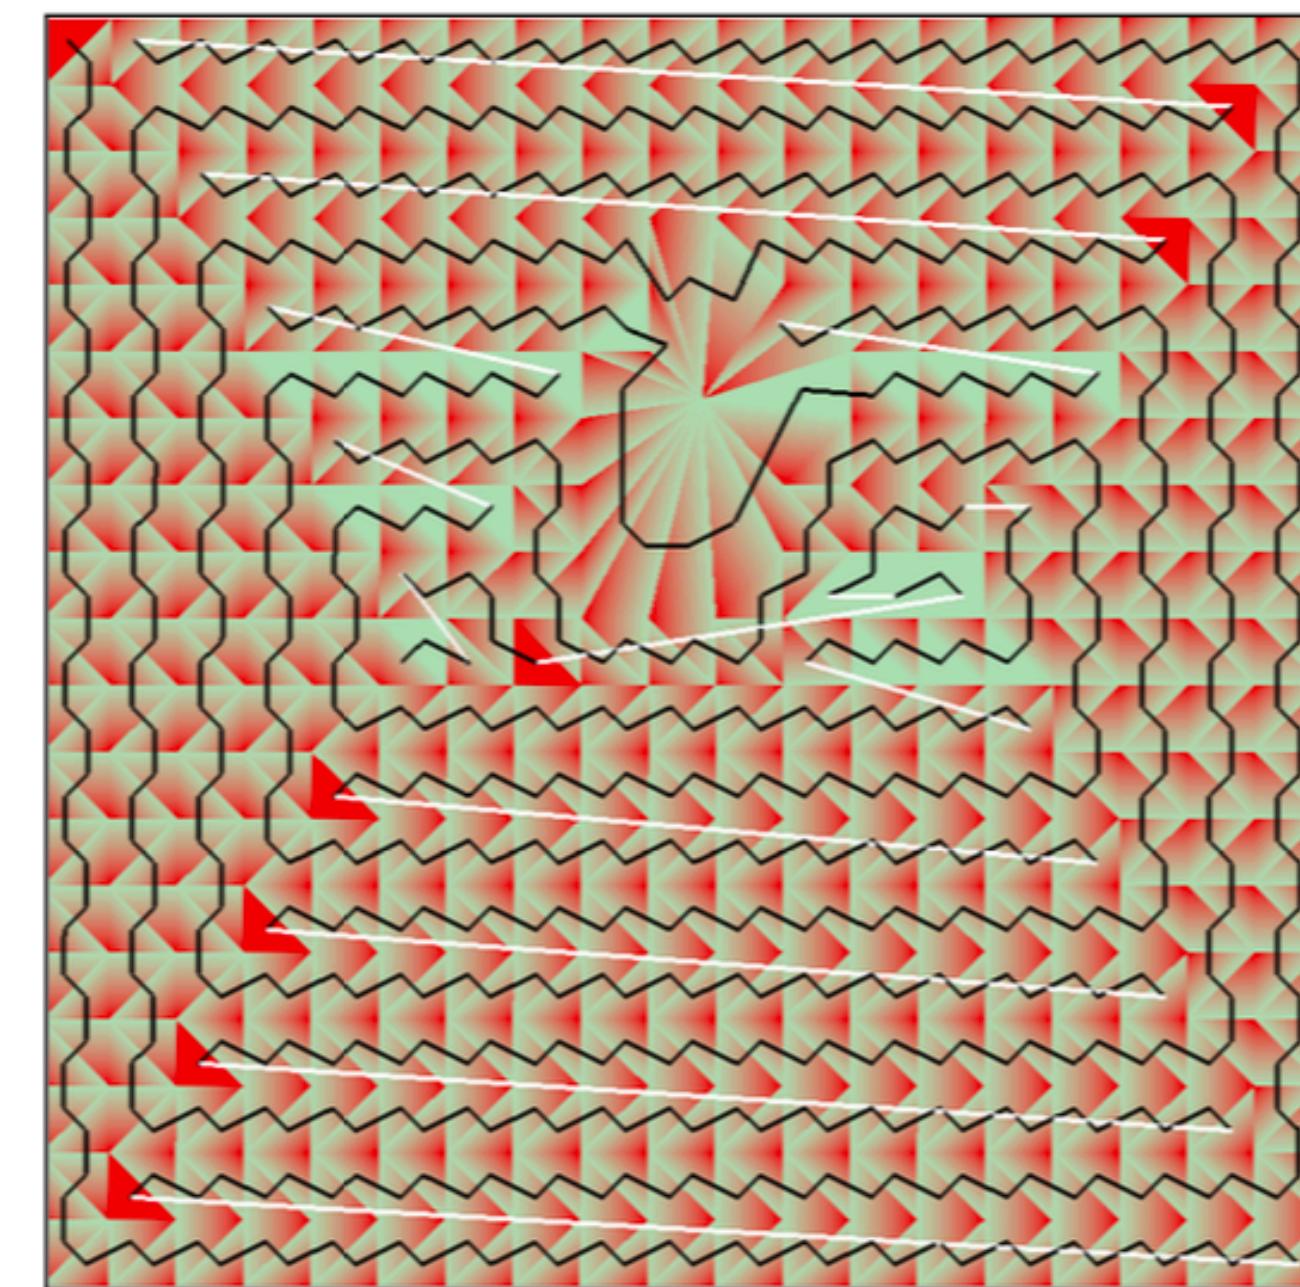
## **Newer systems are built with post-transform vertex caches**

- cache the results of the vertex processing stage
- cache hits can save substantial computation

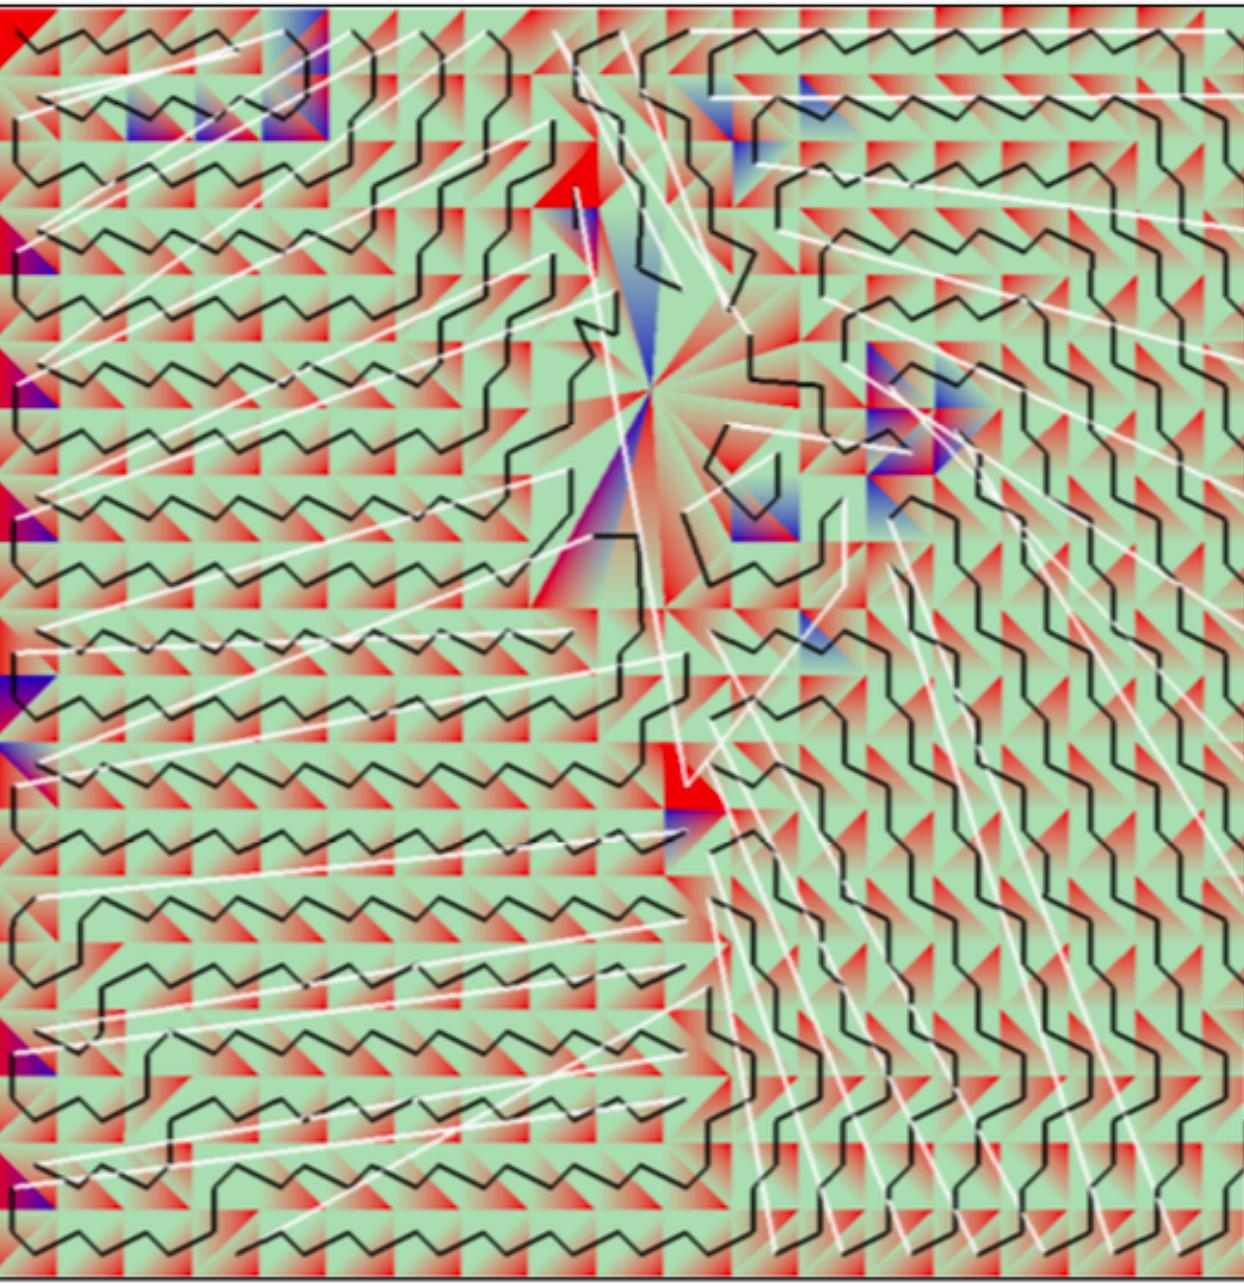
## **As with other applications of caches, now order of data access matters**



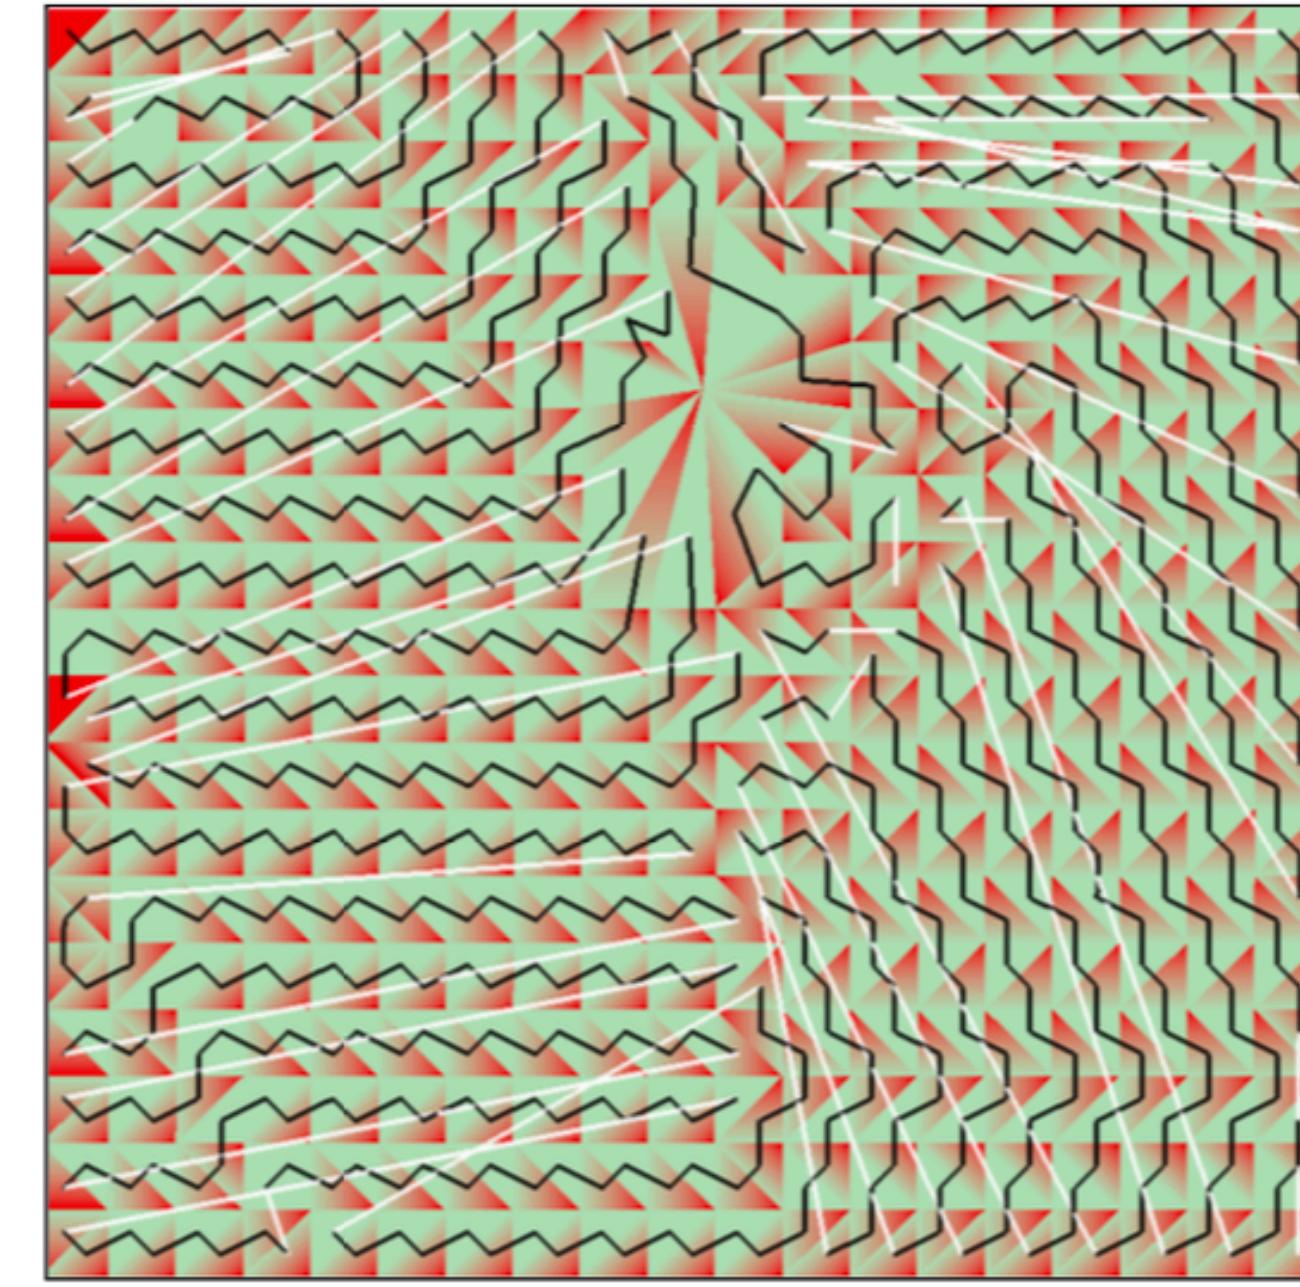
(a) Original mesh (704 faces)



(b) Traditional strips ( $r = 0.99$ ;  $b = 1.29$ ;  $\mathcal{C} = 34.3$ )

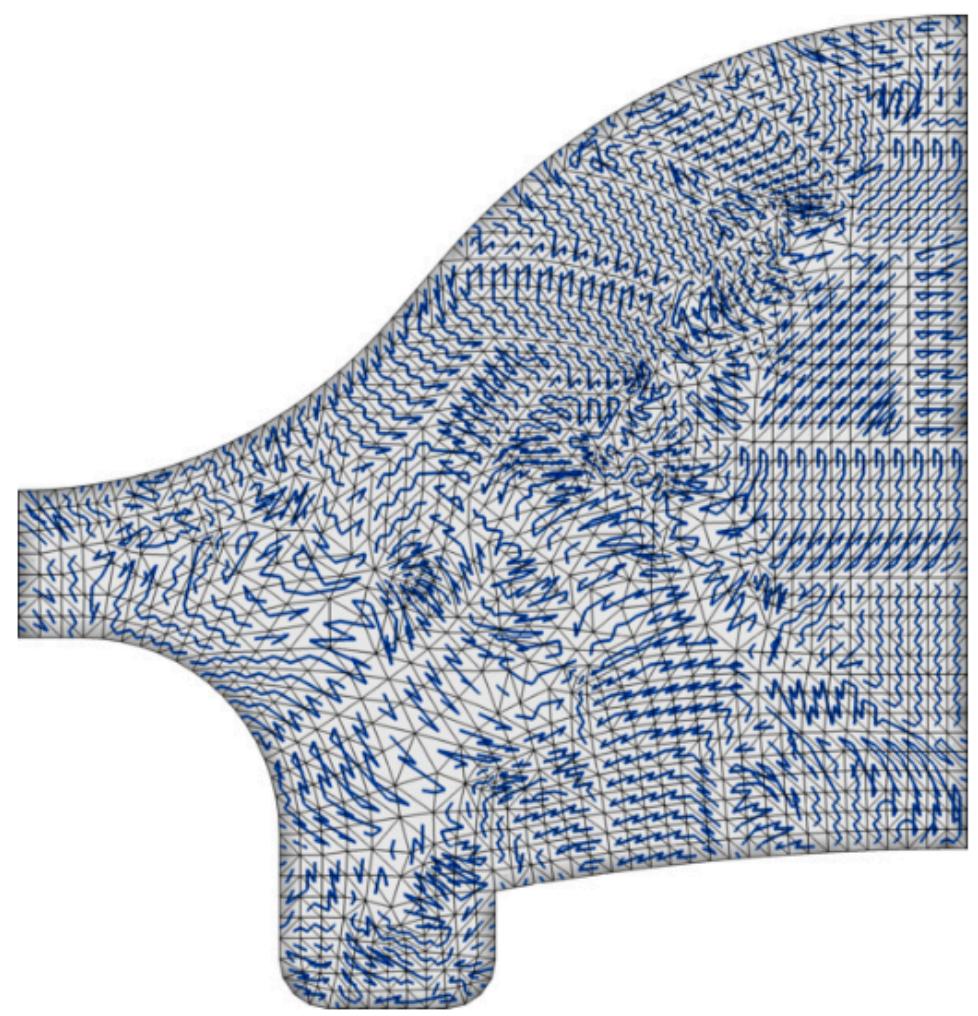


(c) Greedy strip-growing ( $r = 0.62$ ;  $b = 1.28$ ;  $\mathcal{C} = 22.3$ )

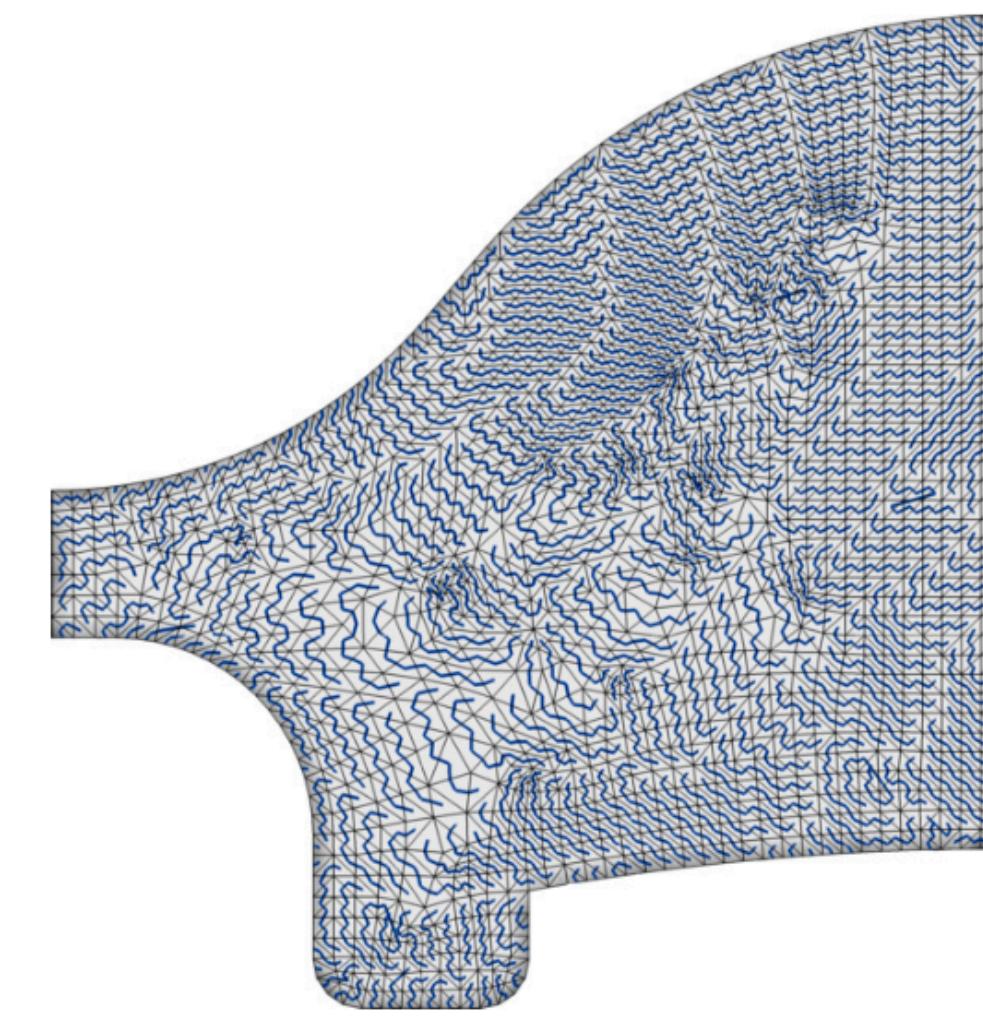


(d) Local optimization ( $r = 0.60$ ;  $b = 1.32$ ;  $\mathcal{C} = 21.7$ )

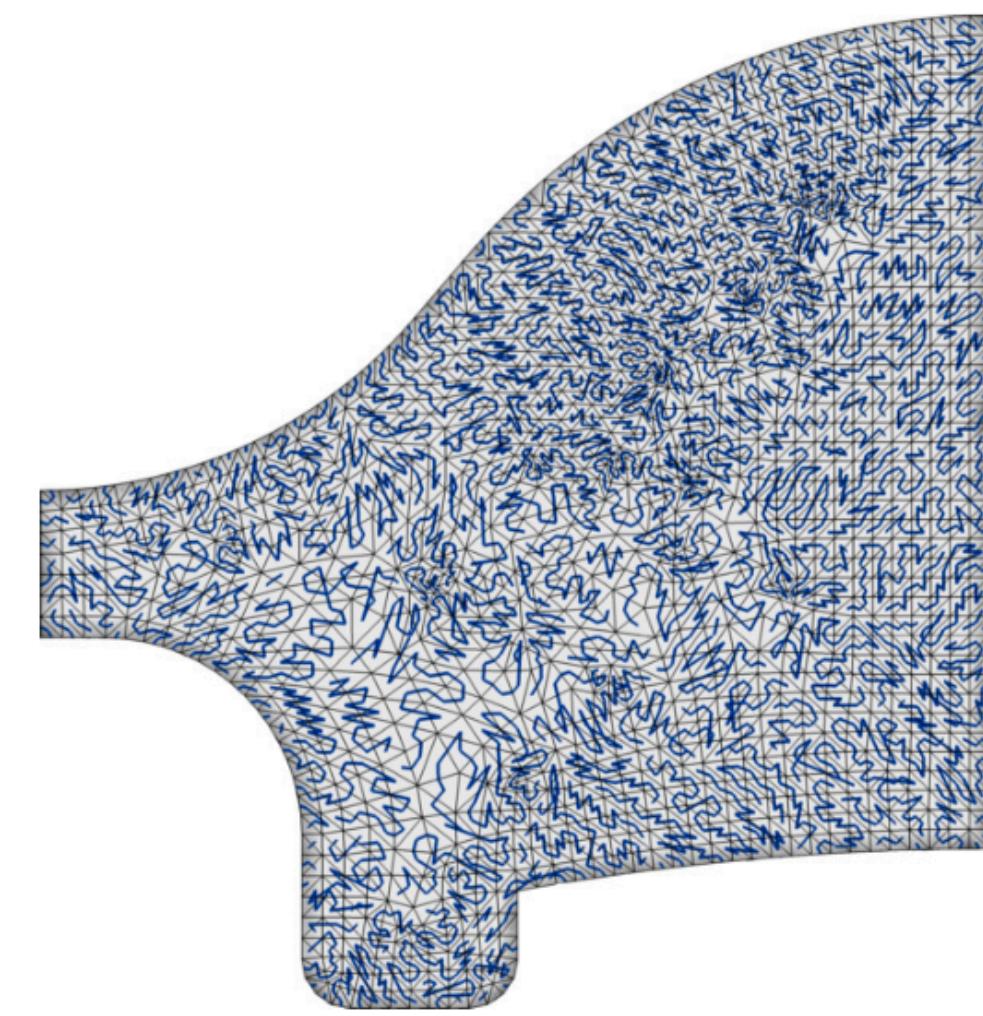
$r$  is the average  
cache miss rate



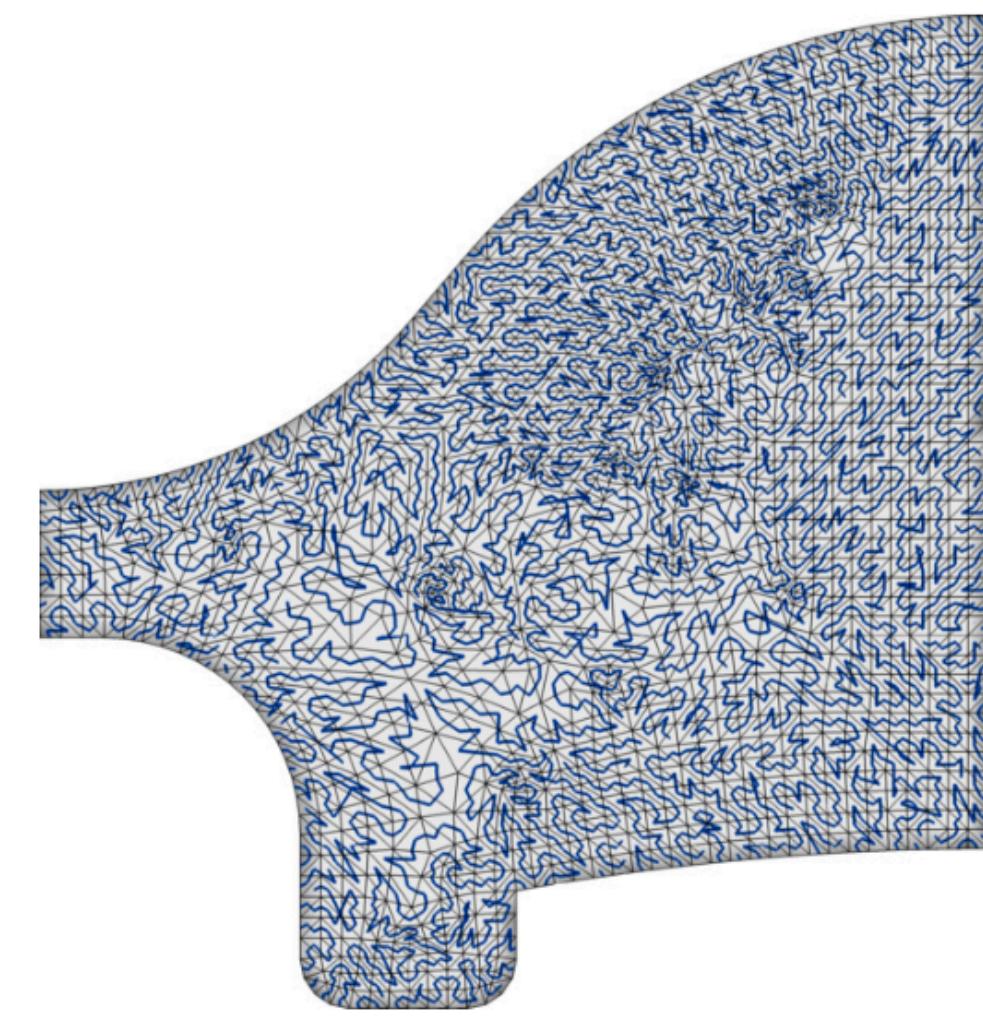
(a) *K-Cache-Reorder*  
Lin and Yu [2006]



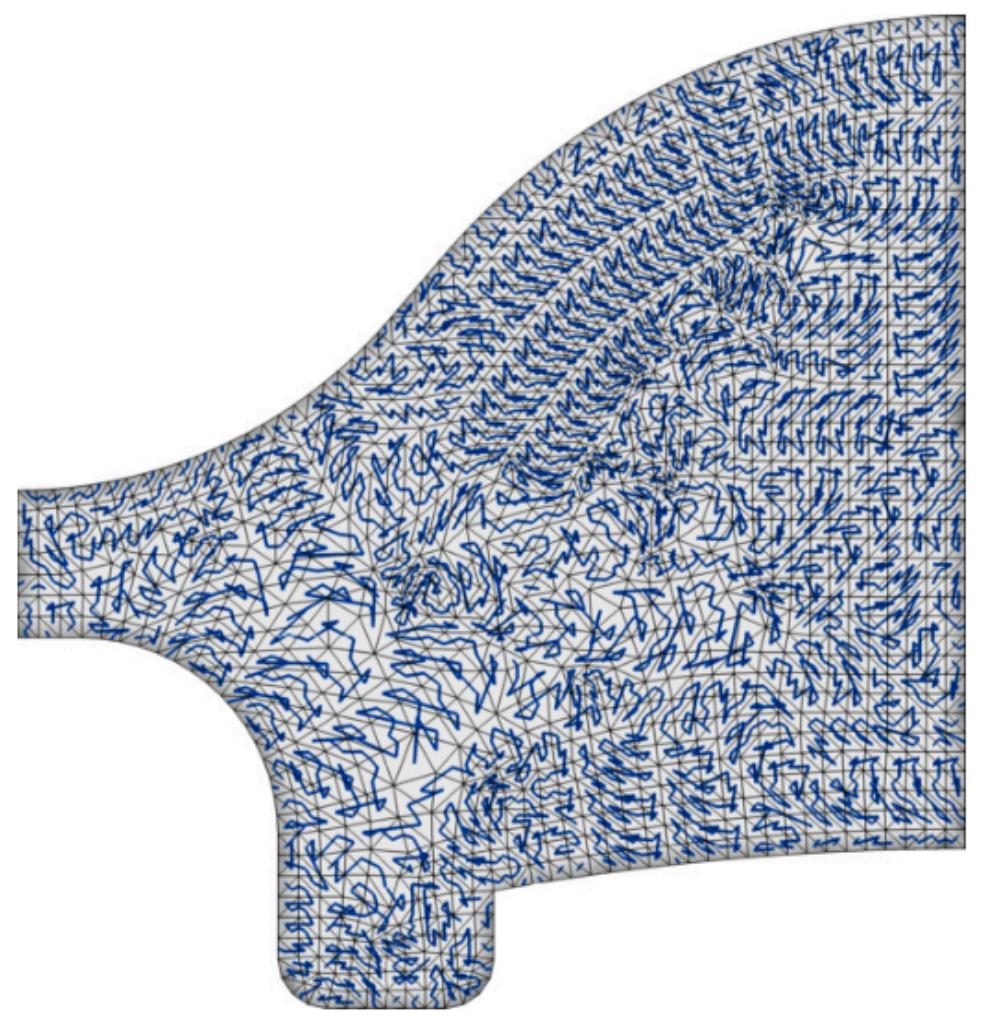
(b) *D3DXMesh*  
based on [Hoppe 1999]



(c) *OpenCCL*  
Yoon and Lindstrom [2006]



(d) *dfsrendseq*  
Bogomjakov et al. [2001]



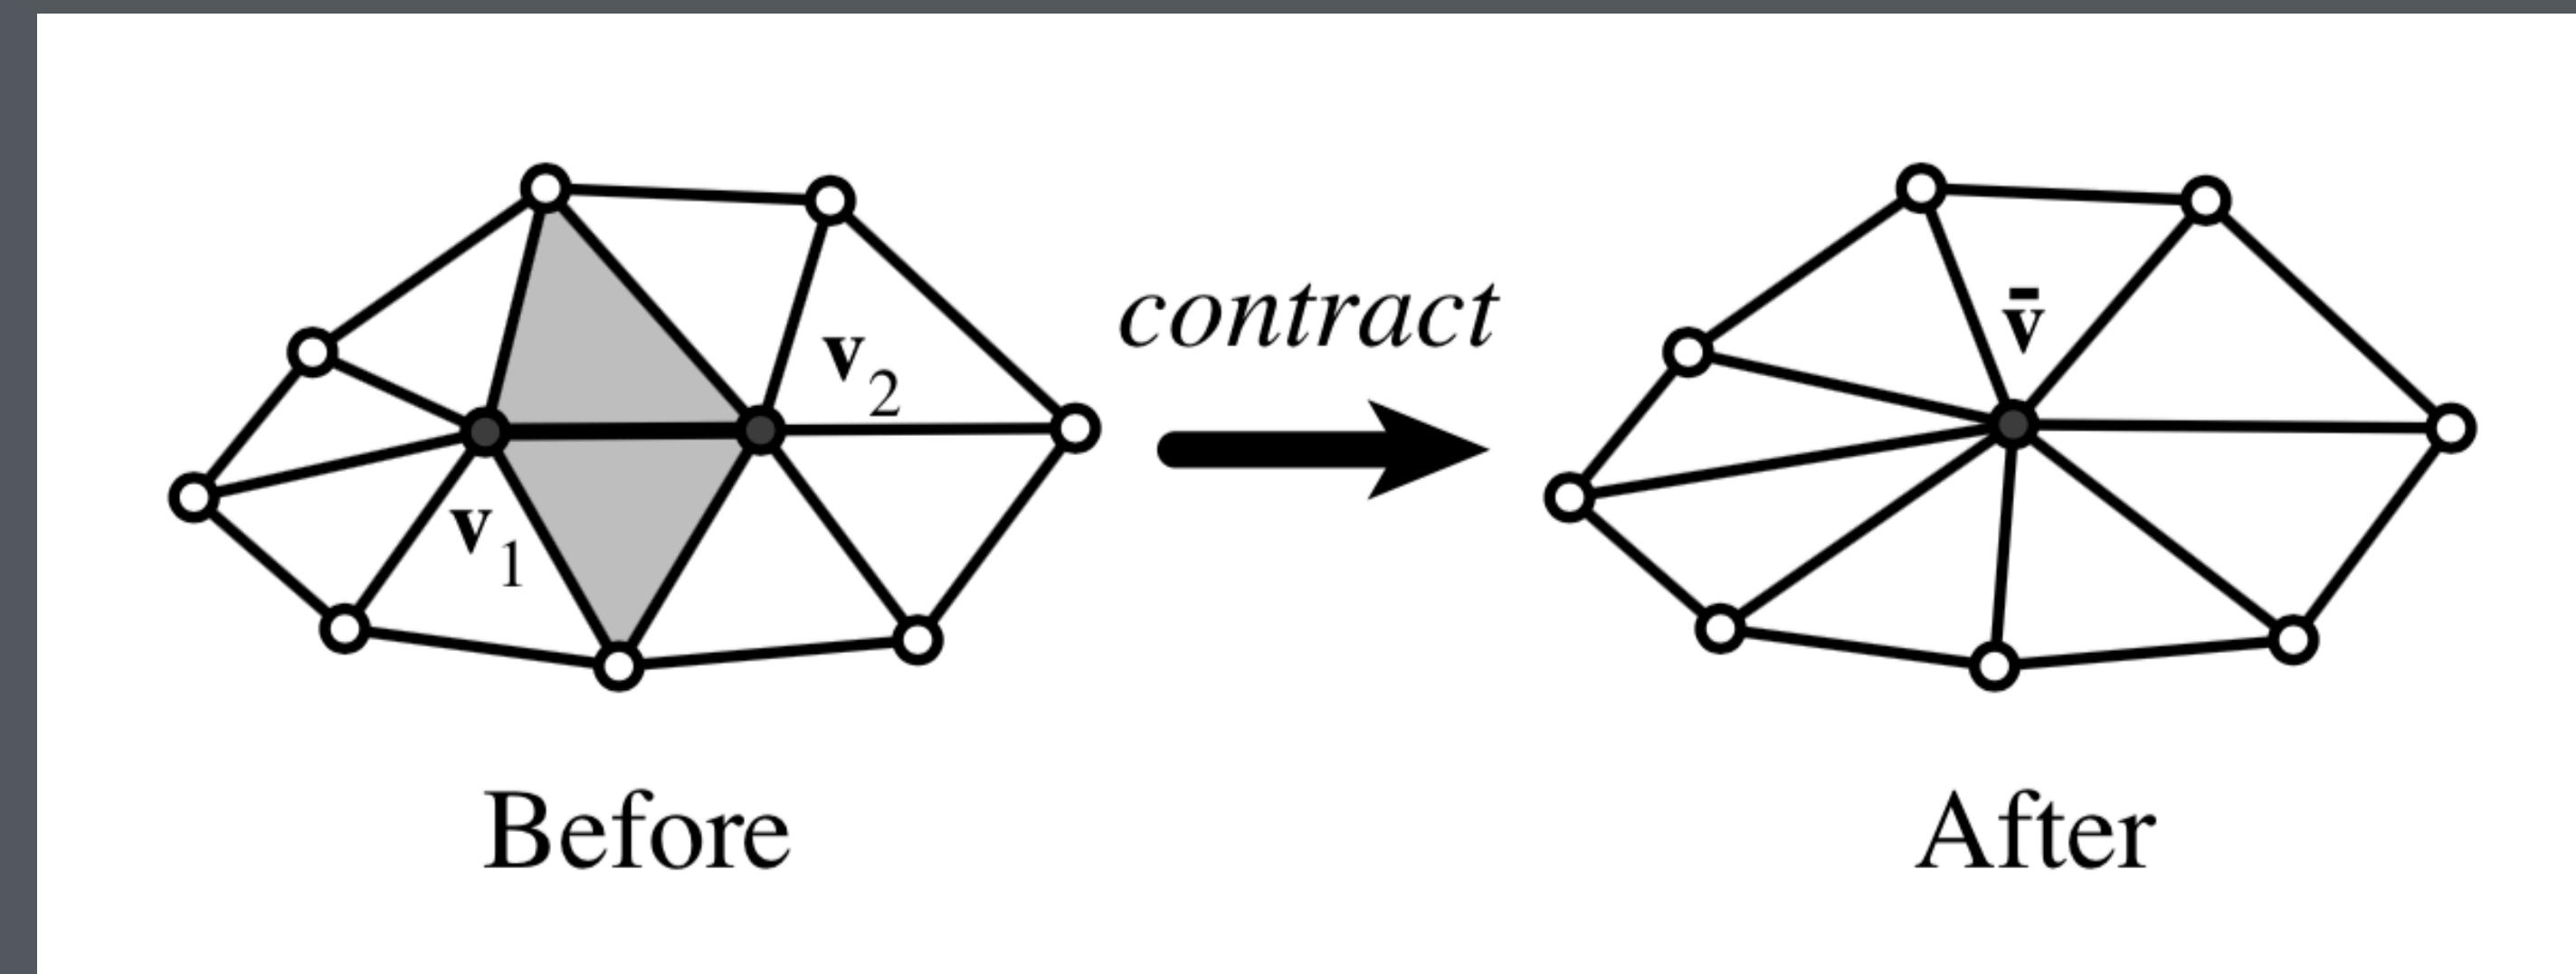
(e) *Our work*

# Mesh simplification

## Many ways to simplify meshes

- remove chunks, retriangulate hole
- quantize vertices to centers of voxels

Particularly simple and effective is edge collapses, or edge contractions:



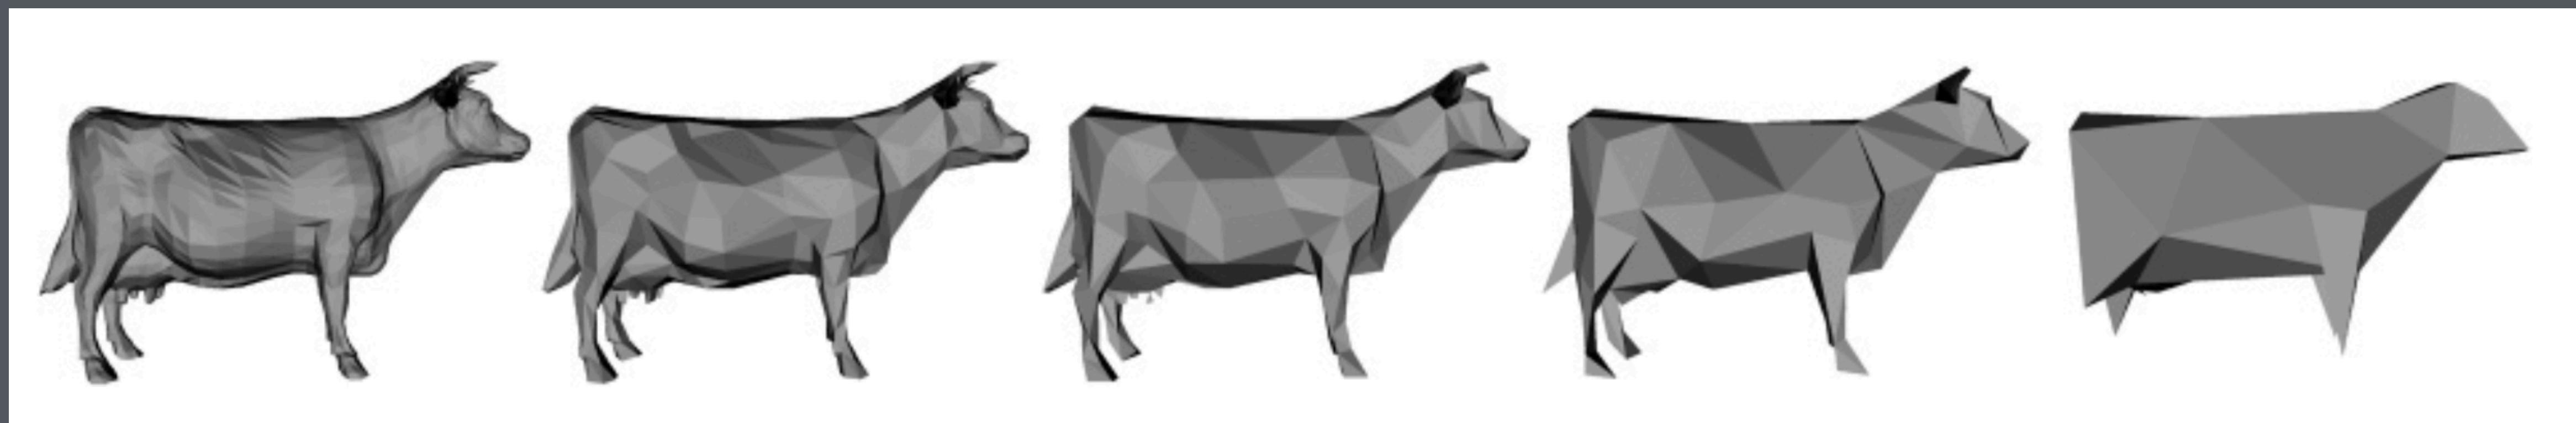
# Quadric Error Metric

**Edge-collapse simplification produces a sequence of meshes**

- each mesh has one fewer face
- each is derived from the previous by a single edge collapse

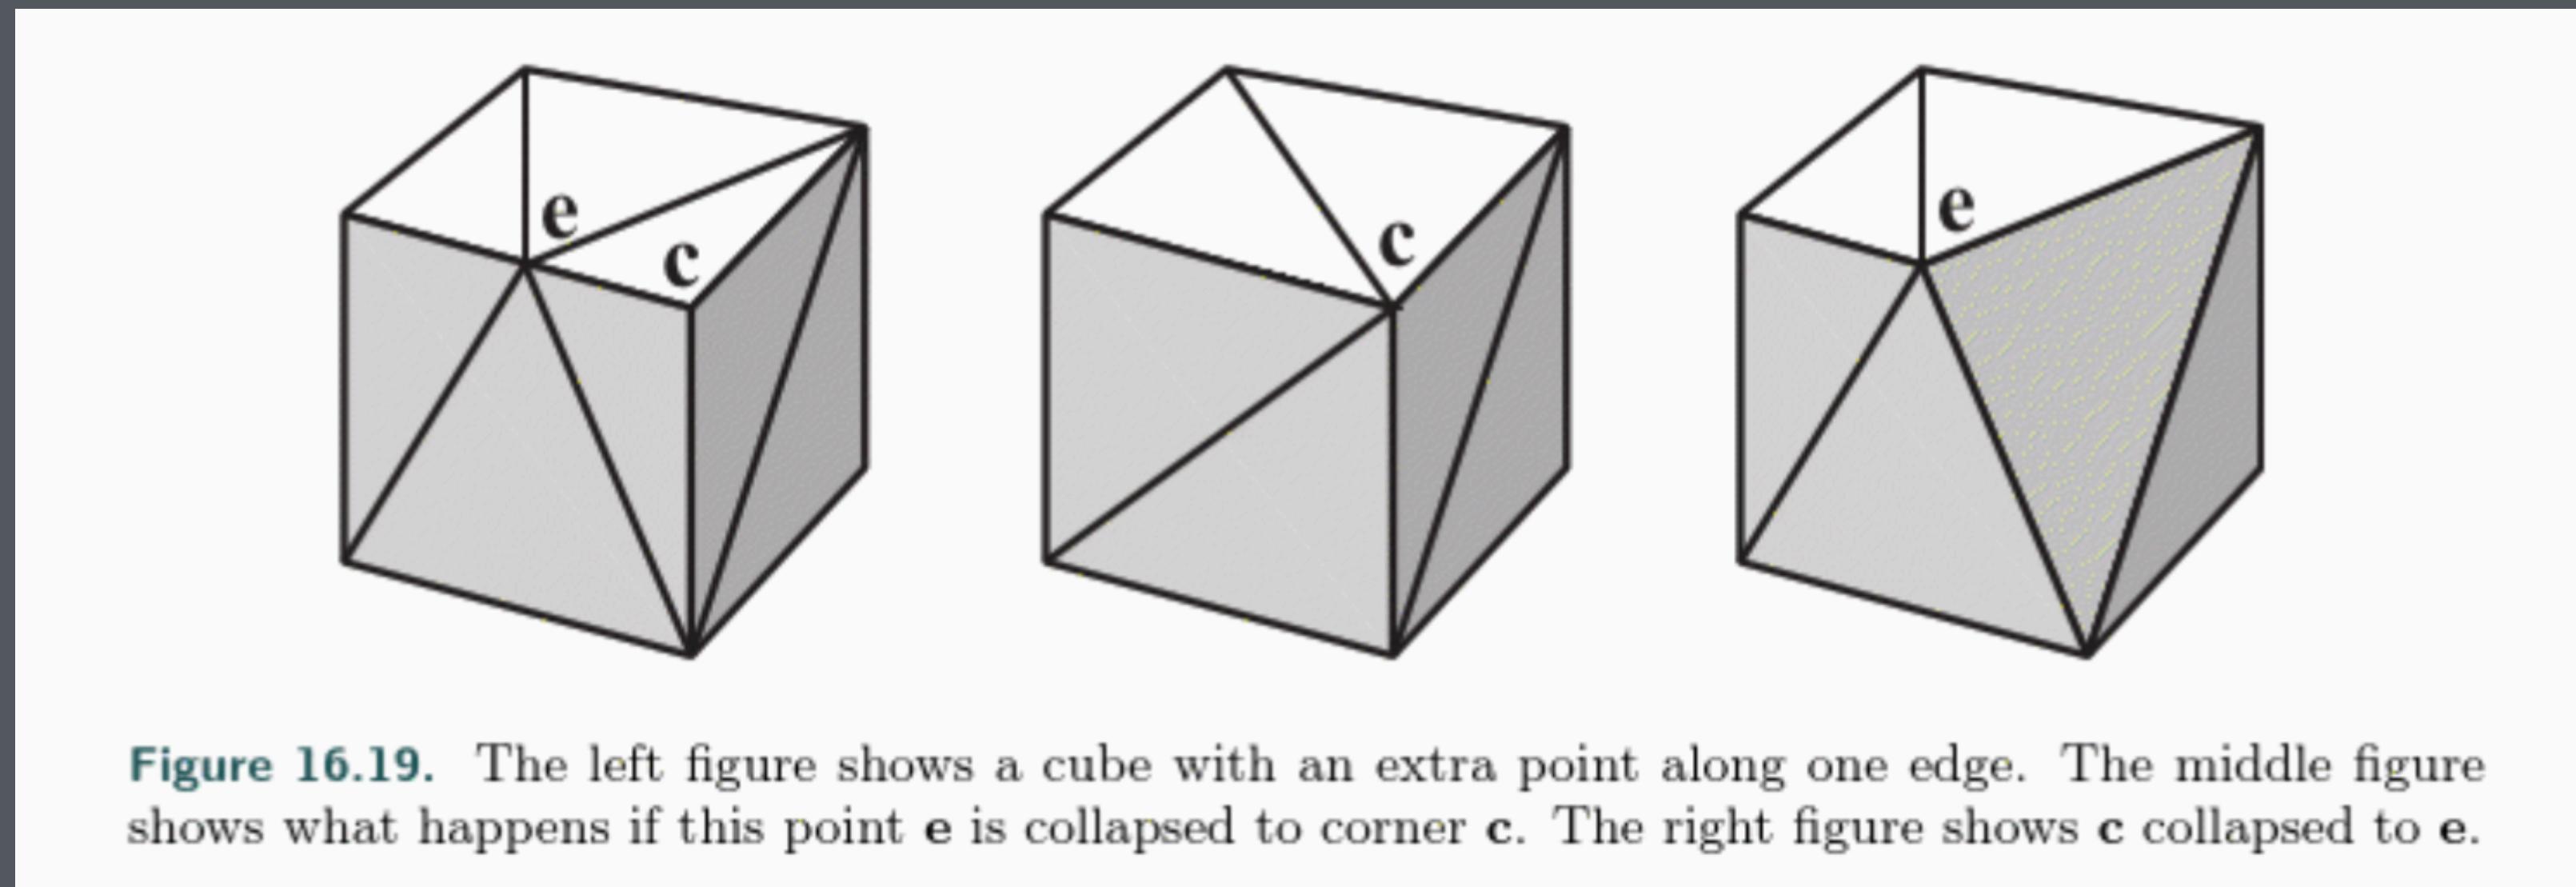
**Key question: where to put the vertex after the collapse?**

- at first vertex? at second? at midpoint?
- can choose location as the solution to an optimization



# Where to put the new vertex?

It depends on the mesh geometry:



- one way to formalize: the new vertex should be close to the planes of the triangles around it before the edge collapse

# Garland & Heckbert QEM

## A particularly convenient error metric: sum of squared distances to planes

- each plane has an equation, can be represented as a 4-vector  $(a, b, c, d)$  with  $(a, b, c)$  components normalized
- distance of a vertex  $\mathbf{v}$  from the plane  $\mathbf{p}$  is then the inner product  $\mathbf{p}^T \mathbf{v}$
- squared distance from plane is in the form  $\mathbf{v}^T \mathbf{M} \mathbf{v}$  for a  $4 \times 4$   $\mathbf{M}$  (a quadric)

$$\begin{aligned}\Delta(\mathbf{v}) &= \sum_{\mathbf{p} \in \text{planes}(\mathbf{v})} (\mathbf{v}^T \mathbf{p})(\mathbf{p}^T \mathbf{v}) \\ &= \sum_{\mathbf{p} \in \text{planes}(\mathbf{v})} \mathbf{v}^T (\mathbf{p} \mathbf{p}^T) \mathbf{v} \\ &= \mathbf{v}^T \left( \sum_{\mathbf{p} \in \text{planes}(\mathbf{v})} \mathbf{K}_p \right) \mathbf{v}\end{aligned}$$

- and better yet, the sum-squared distance from several planes is still in the form  $\mathbf{v}^T \mathbf{Q} \mathbf{v}$

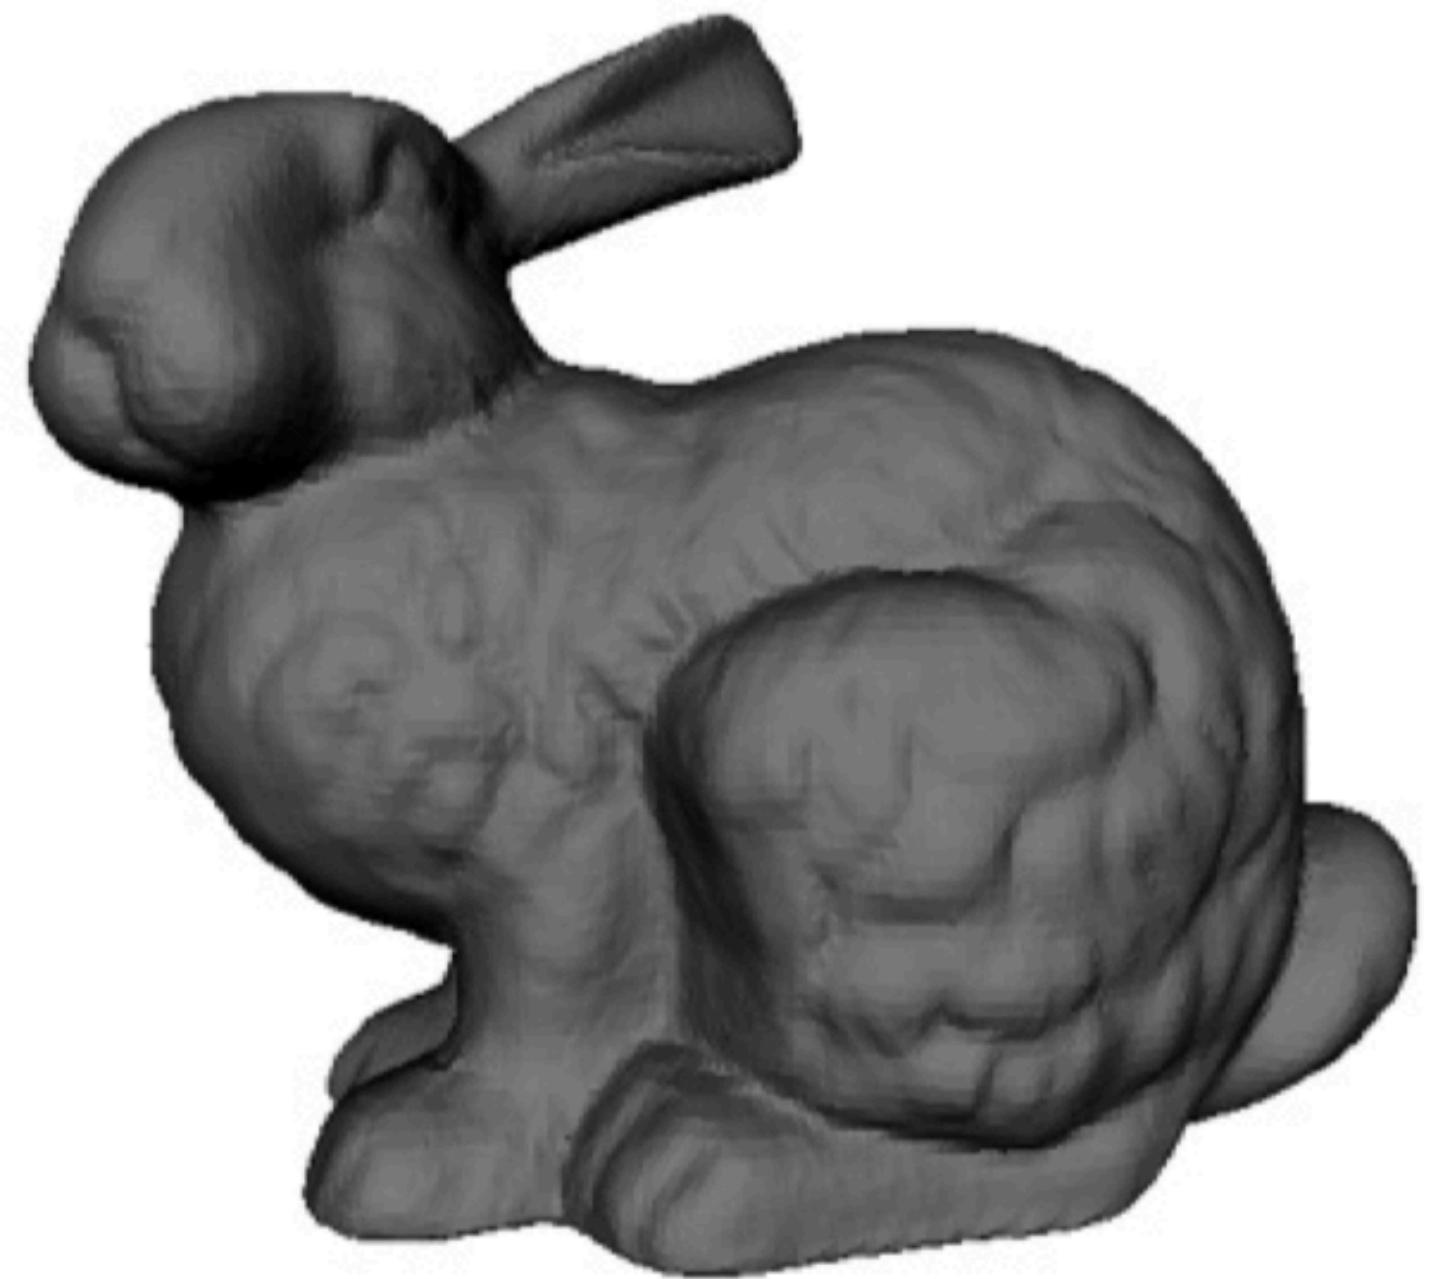
# QEM simplification

## With the error in the form of a quadric per vertex:

- the matrix is easy to compute from the surrounding triangles
- the error is easy to optimize. Given  $\mathbf{Q}_1$  and  $\mathbf{Q}_2$  belonging to a pair of vertices  $\mathbf{v}_1$  and  $\mathbf{v}_2$ , we simply sum the errors of the two vertices:

$$\begin{aligned}\Delta(\mathbf{v}) &= \Delta_1(\mathbf{v}) + \Delta_2(\mathbf{v}) \\ &= \mathbf{v}^T \mathbf{Q}_1 \mathbf{v} + \mathbf{v}^T \mathbf{Q}_2 \mathbf{v} \\ &= \mathbf{v}^T (\mathbf{Q}_1 + \mathbf{Q}_2) \mathbf{v}\end{aligned}$$

- minimizing this error is a  $4 \times 4$  linear system—very fast
- algorithm
  - 0. compute  $\mathbf{Q}$ s for all vertices, compute errors for all potential edge collapses.
  - 1. use priority queue to find smallest-error edge. Collapse it; update the neighboring  $\mathbf{Q}$ s.
  - 2. repeat until mesh is small enough!



69k faces



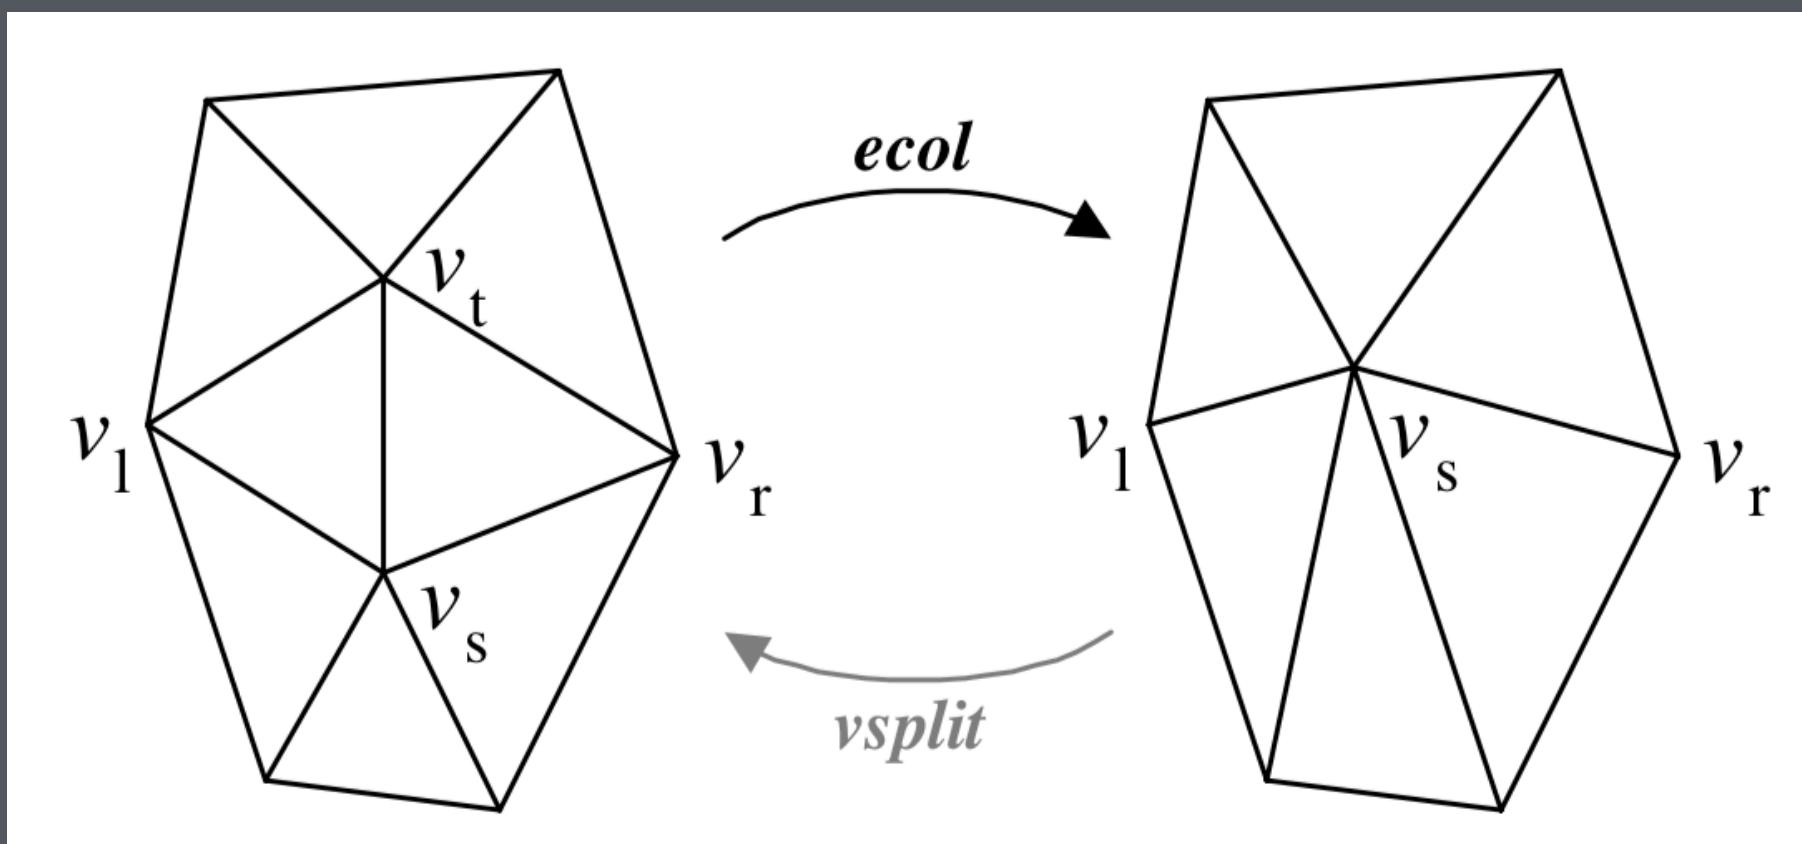
1k faces



surfaces of constant  
cost for reducing to  
999 faces

# Continuous level-of-detail: Progressive Meshes

**Key observation: edge collapse is invertible**



- just need to store (offsets to) the locations of the two new vertices

**Thus a sequence of edge collapses, reversed, is a representation for a mesh**

$$(\hat{M} = M^n) \xrightarrow{\text{ecol}_{n-1}} \dots \xrightarrow{\text{ecol}_1} M^1 \xrightarrow{\text{ecol}_0} M^0 .$$

$$M^0 \xrightarrow{\text{vsplit}_0} M^1 \xrightarrow{\text{vsplit}_1} \dots \xrightarrow{\text{vsplit}_{n-1}} (M^n = \hat{M})$$

# Progressive Meshes

## **Store full representation, load various levels of detail**

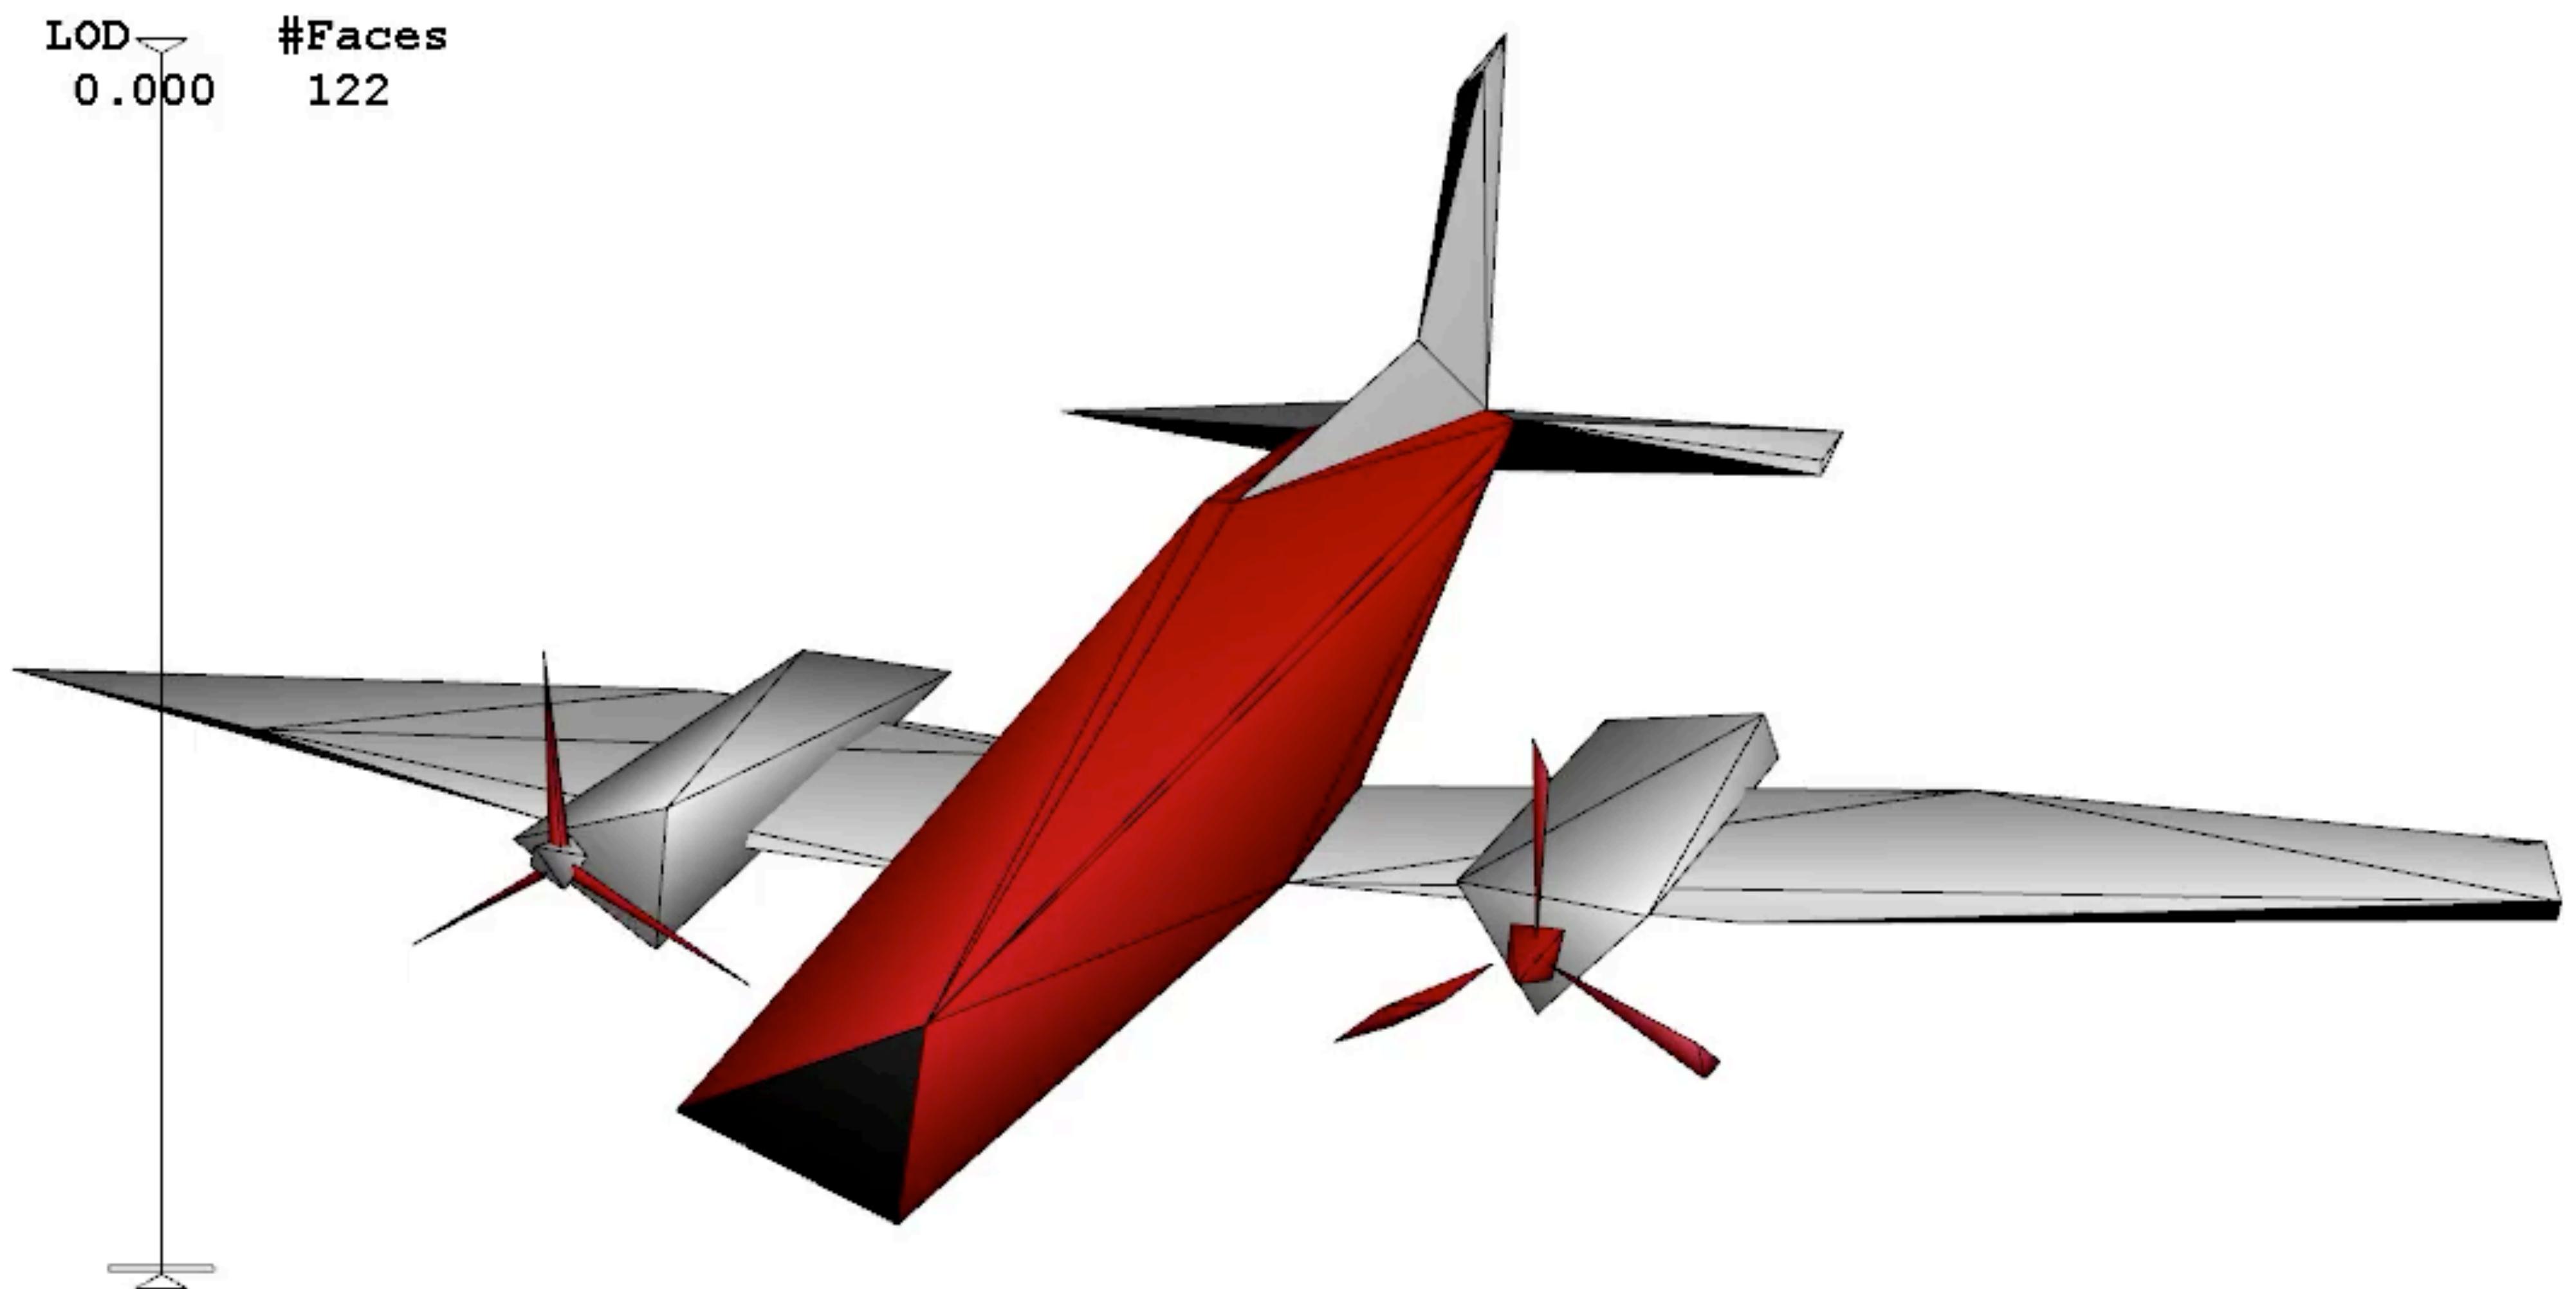
- just load or transmit a prefix of the list of edge splits
- can change level of detail smoothly depending on size/distance/salience/etc.

## **Can interpolate (“geomorphs”)**

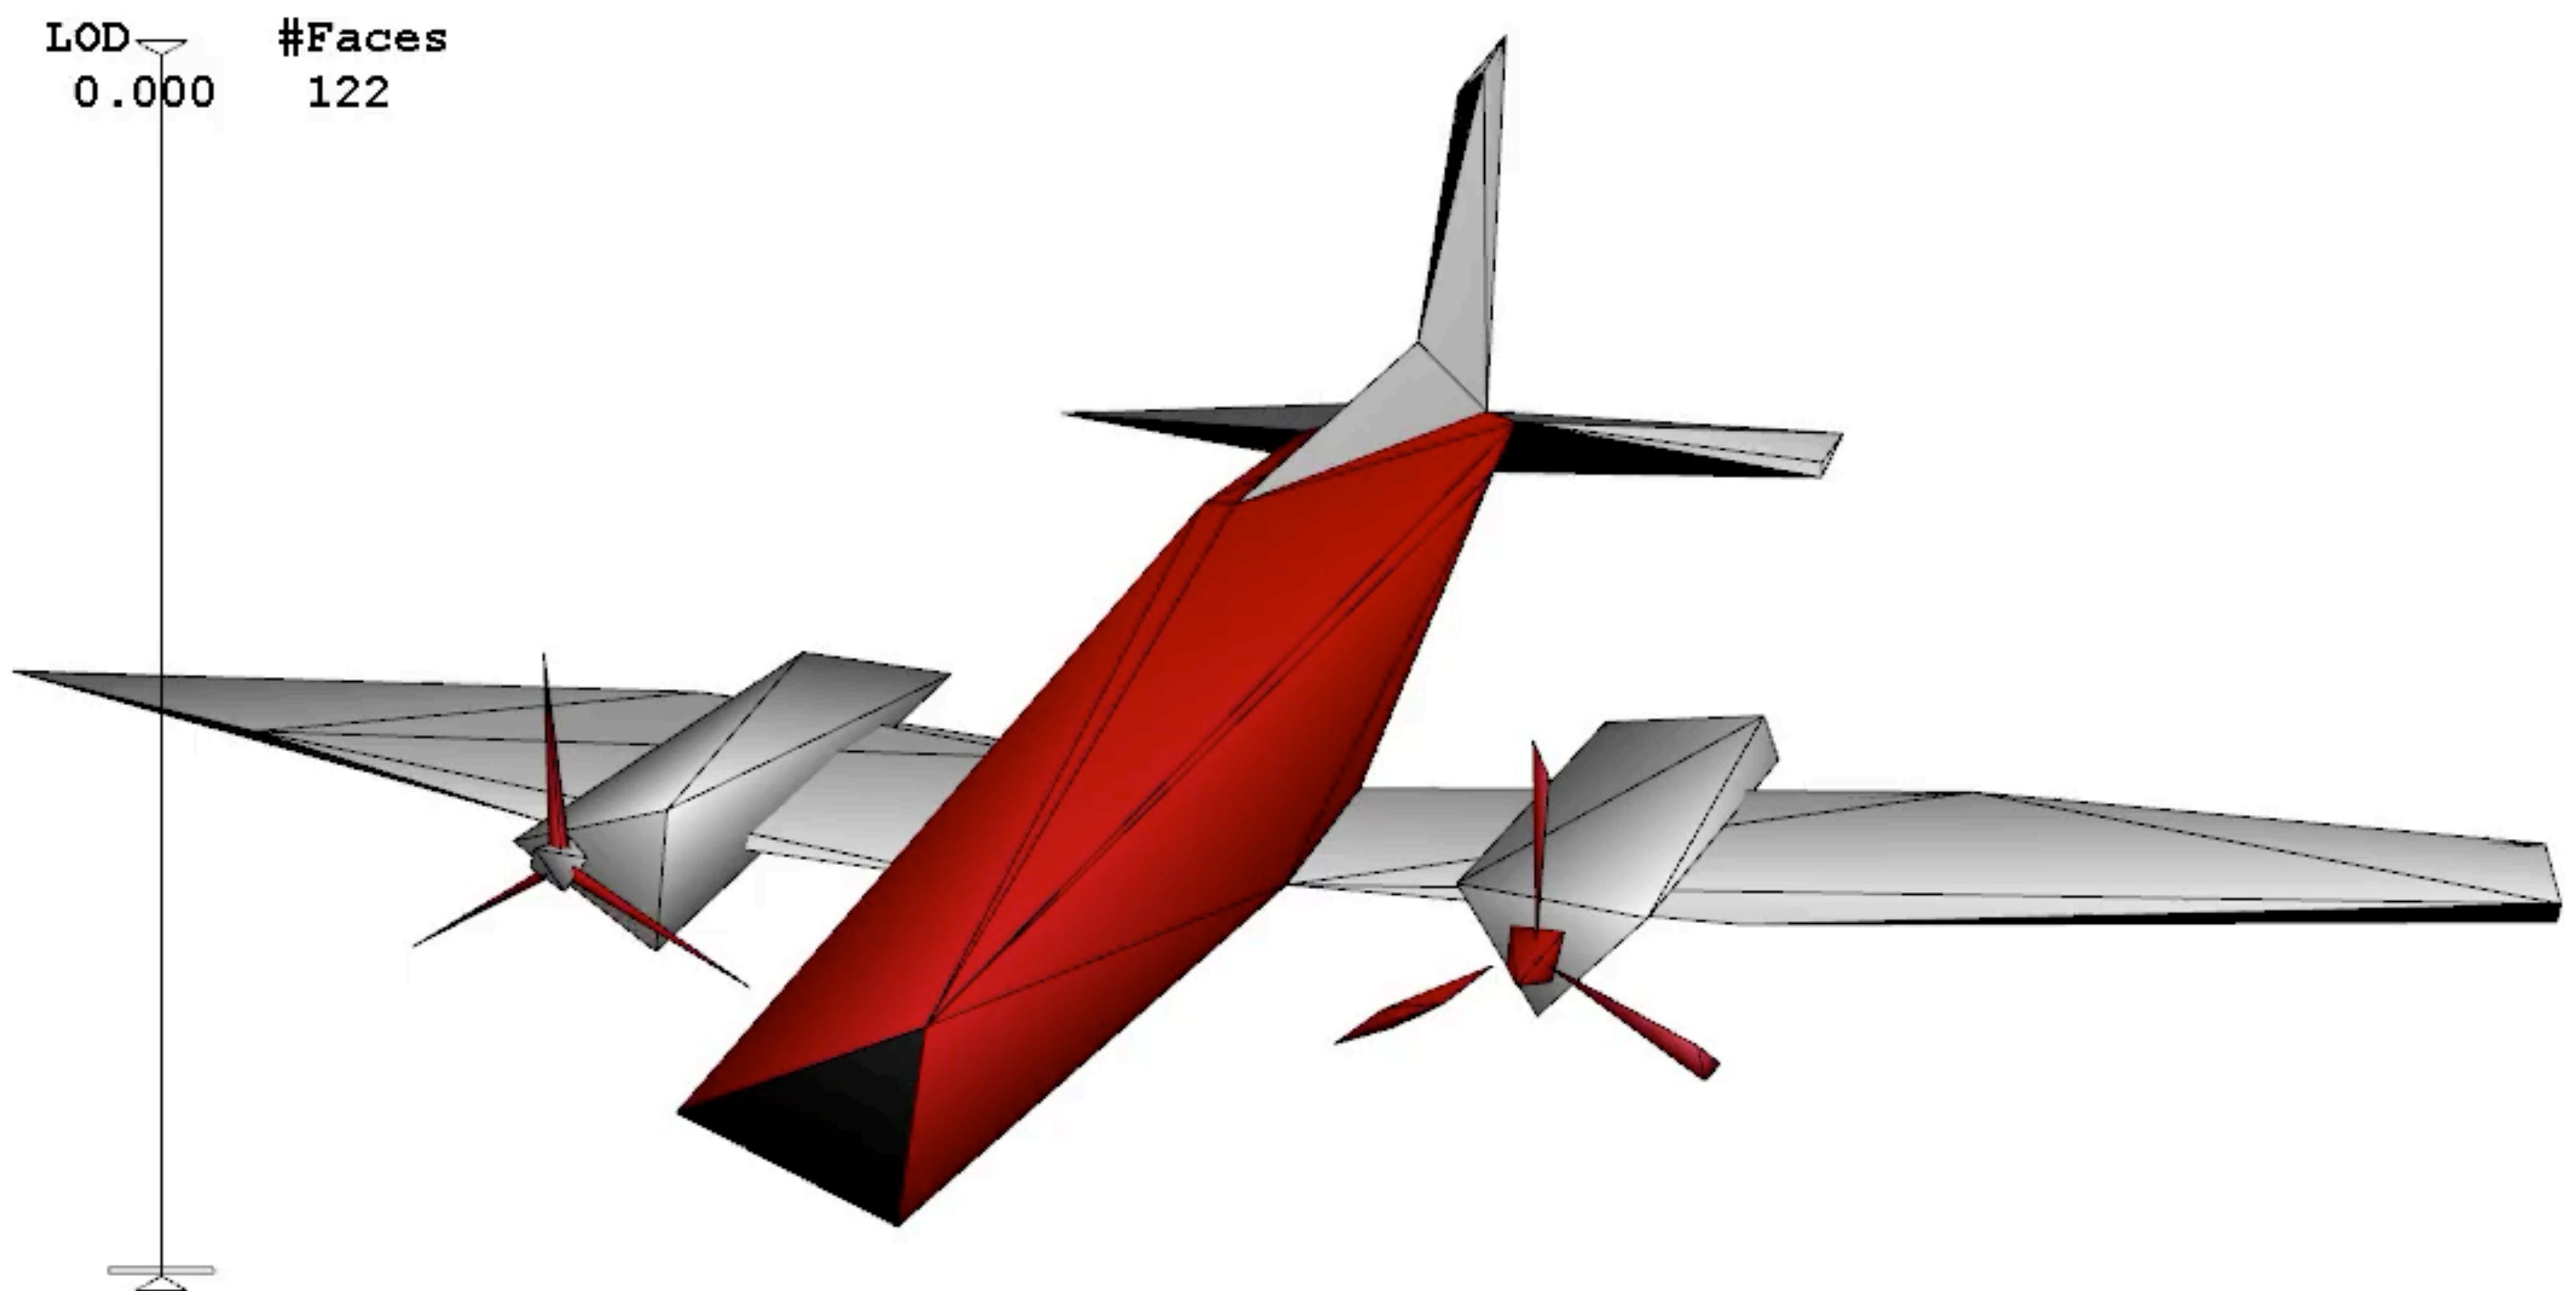
- sudden edge splits/collapses are jarring
- interpolate new vertices from merged position to new positions
- leads to truly continuous LoD

## **Extra details (of QEM and PM)**

- boundaries, creases—want to preserve them
- merging of small pieces—otherwise can’t simplify enough
- maintenance of additional attributes—throw them in the metric too



[Hoppe 1996 “Progressive Meshes”]



[Hoppe 1996 “Progressive Meshes”]